

3. Data Structures and Tolerant Retrieval

Prof. Dr. Goran Glavaš

Center for AI and Data Science (CAIDAS)
Fakultät für Mathematik und Informatik
Universität Würzburg



Creative Commons Attribution-NonCommercial-ShareAlike 4.0 International

After this lecture, you'll...

2

- Know what data structures are used for implementing inverted index
- Understand the pros and cons of hash tables and trees
- Know how to handle wildcard queries
- Be familiar with methods for handling spelling errors and typos in IR

Outline

3

- **Recap of Lecture #2**
- Data structures for inverted index
- Wild-card queries
- Spelling correction

Recap of the previous lecture

4

- Boolean retrieval
 - **Q:** How are queries represented in Boolean retrieval?
 - **Q:** How are documents represented for Boolean retrieval?
 - **Q:** How do we find relevant documents for a given query?
- Inverted index and finding relevant documents
 - **Q:** What is inverted index and what does it consist of?
 - **Q:** What are posting lists?
 - **Q:** How to merge posting lists?
 - **Q:** What is the computational complexity of the merge algorithm?
 - **Q:** What are skip pointers and what is their purpose?
- Phrase and proximity queries
 - **Q:** What is a biword index and what are its shortcomings?
 - **Q:** What is a positional index?
 - **Q:** How do we use positional index to answer phrase and proximity queries?

Recap of the previous lecture

5

- **Inverted index** is a data structure for computationally efficient retrieval
- Inverted index contains a list of references to documents for all index terms
 - For each term *t* we store the list of all documents that contain *t*
 - Documents are represented with their identifier numbers (ordinal, starting from 1)

„Frodo” -> [1, 2, 7, 19, 174, 210, 331, 2046]

„Sam” -> [2, 3, 4, 7, 11, 94, 210, 1137]

„blue” -> [2, 3, 24, 2001]

- The list of documents that contains a term is called a **posting list** (or just **posting**)
- **Q:** Postings are always sorted. Why?

Recap of the previous lecture

6

- So far, we learned how to handle
 - Regular Boolean queries
 - Standard merge algorithm over posting lists
 - Multi-term queries – optimizing according to lengths of posting lists
 - Phrase queries
 - Biword index
 - Positional index
 - Proximity queries
 - Positional index
- **Today we'll examine**
 - Data structures for implementing the inverted index
 - How to handle wild-card queries and spelling errors

Outline

7

- Recap of Lecture #2
- Data structures for inverted index
- Wild-card queries
- Spelling correction

Data structures for inverted index

8

- Conceptually, an inverted index is a **dictionary**
 - Vocabulary terms (i.e., index terms) are keys
 - Posting lists are values

„Frodo” -> [1, 2, 7, 19, 174, 210, 331, 2046]

„Sam” -> [2, 3, 4, 7, 11, 94, 210, 1137]

„blue” -> [2, 3, 24, 2001]

- But the exact implementation is undefined
 - What **data structures** to use?
 - Where exactly to store different pieces of information – document frequencies, pointers to posting lists, skip pointers, token positions, ...?

Data structures for inverted index

9

- A naïve dictionary – an **array/list** of structures

Term	Doc. freq.	Pointer
a	656 265	→
aachen	65	→
blue	10 321	→
...
frodo	221	→

- Each element of the array is a structure consisting of:
 - The term itself
 - The number of documents in the collection in which the term appears
 - A pointer to the posting list of the term
- Structure size: char[N], int, pointer (int/long)

- **Q:** How to efficiently store the inverted index / dictionary in memory?
- **Q:** How to quickly look up elements at query time?

Data structures for inverted index

10

- Two main choices for implementing the inverted index dictionary
 - Hash tables
 - Trees
- Both are regularly used in IR systems
- Both have advantages and shortcomings

Inverted index dictionary as a hash table


11

- **Hash table** is a common data structure for implementing dictionaries, i.e., a structure that maps keys to values as **associative arrays**
- Hash tables rely on **hash functions** – functions that for a given input value (i.e., **key**) computes the index in the array where the **value** is stored
 - Ideally, each key would be assigned a different index
 - Most hash functions are imperfect – they may compute the same value for several different keys – this is called a **collision**
 - **Q:** how to account for collisions?
- Each vocabulary term is „hashed” into an integer value
 - $hf(„Frodo”) = 1$, $hf(„Sam”) = 19$, $hf(„blue”) = 204$

Inverted index dictionary as a hash table

12

- $hf(„Frodo”) = 1, hf(„Sam”) = 19, hf(„blue”) = 204, hf(„orc”) = 1$
- **Associative array**



- If the hash function maps the key to the bucket with more than one entry, then the **linear search through the bucket** is performed

Inverted index dictionary as a hash table

13

- The main advantage of hash table is **fast lookup**
 - **Q:** What is the complexity of the lookup?
 - **A:** $O(1)$
- **Shortcomings:**
 - Hash functions are sensitive to minor differences in strings
 - Close strings not assigned same or close buckets
 - E.g., $hf(„judgment”) = 12$, $hf(„judgement”) = 354$
 - As such, they **do not support prefix search**
 - Important for tolerant retrieval
 - Constant vocabulary growth means **occasional rehashing for all terms**
 - **Q:** Why do we need to rehash if the vocabulary grows?

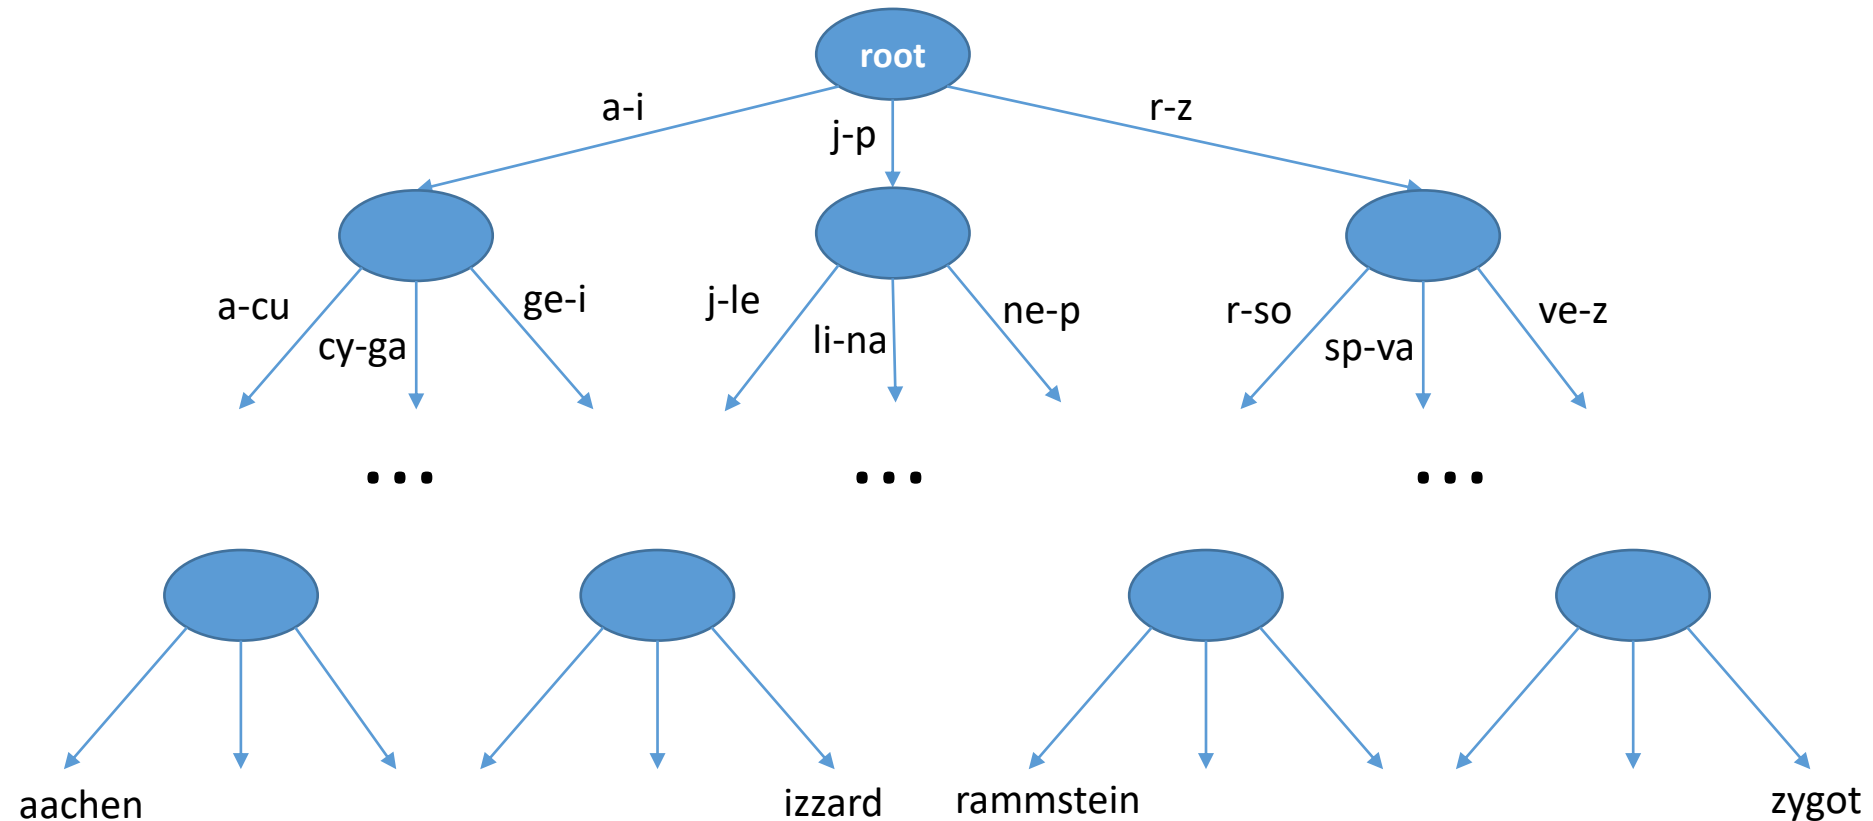
Inverted index dictionary as a tree

14

- **Trees** divide the vocabulary in the hierarchical form
 - Each node in the tree captures the **subset of the vocabulary**
 - Nodes closer to the root represent larger vocabulary subsets
 - Nodes closer to leaves encompass narrower subsets
 - Actual **vocabulary terms** are found in **leaf nodes** of the tree
 - The division of vocabulary is usually **alphabetical**
- Trees should be created in a **balanced fashion**
 - Each node in the tree should have approximately the same number of children
 - Subtrees of nodes at the same depth should have approx. the same number of leaves

Inverted index dictionary as a tree

15



Inverted index dictionary as a tree

16

- **Q:** What is the lookup complexity for a balanced tree with a node degree N which stores vocabulary containing $|V|$?
 - **A:** Lookup complexity is equal to the depth of the tree, so the complexity is $O(\log_N |V|)$
- The central design decision is the **degree** of the nodes in an index tree, i.e., the number of child nodes a parent node should have
 - Large node degree N
 - Shallow trees, but a large number of children to go through linearly
 - Small node degree N
 - Small number of children to linearly search, but deep trees
- Advantage
 - Can handle **prefix search**
- Shortcoming
 - Lookup complexity ($O(\log_N |V|)$) bigger than for hash tables ($O(1)$)

Outline

17

- Recap of Lecture #2
- Data structures for inverted index
- Wild-card queries
- Spelling correction

Wild-card queries

18

- **Wild-card queries** are queries in which an asterisk sign stands for any sequence of characters
 - Wild-card term (with an asterisk) represents a group of terms and not a single term
- Trailing wild-card queries (aka **prefix queries**, * at the end)
 - E.g., „**mon***” is looking for all documents containing any word beginning with „**mon**”
 - Easy to handle with **B-tree index**: retrieve all words **w** in range **mon** \leq **w** $<$ **moo**
- Leading wild-card queries (* at the beginning)
 - E.g., „***mon**” is looking for all documents containing any word ending with „**mon**”
 - Can be handled with an additional B-tree that indexes vocabulary terms **backwards**
 - Retrieve all words **w** in range: **nom** \leq **w** $<$ **non**
- **Q:** How to handle queries with the wild-card in the middle?
 - Retrieve documents containing any word satisfying the wild-card query „**co*tion**”?

Wild-card queries

19

- Example query

- „co*tion” (we want: *coordination*, *comotion*, *cohabitation*, *connotation*, ...)

Idea:

1. Lookup „co*” in the forward B-tree of the vocabulary
2. Lookup „*tion” in the backward B-tree of the vocabulary
3. Intersect the two obtained term sets

- Unfortunately, this is **too expensive (too slow)** for most real-time IR settings
 - We need to fetch the relevant documents with a single lookup into index
 - We need to **enrich the index** somehow
 - This will increase the index size, but memory is usually less of an issue

Wild-card queries and permuterm index

20

- Idea: index all **character-level permutations** of terms
- **Permuterm index** additionally stores permutations of vocabulary terms
- We add a special „end-of-term” character (\$) and store all permutations:
 - E.g., „comotion” -> „\$comotion”, „n\$comotio”, „on\$comoti”, „ion\$comot”, „tion\$como”, „otion\$com”, „motion\$co”, „omotion\$c”
- **Q:** How to use **permuterm index** for **middle-wild-card queries**?
 - **A:** Permute the wild-card query until you obtain a trailing query (asterisk at the end)
 - E.g., „co*tion” -> „\$co*tion” -> „tion\$co*”
 - We know how to handle trailing wild-card queries – „tion\$co*” can now be handled by a single permutex index tree

Wild-card queries and permuterm index

21

- Queries supported by permuterm index
 - Exact queries: for „X” we look up „\$X”
 - Trailing wild-card query: for „X*” we look up „\$X*”
 - Leading wild-card query: for „*X” we look up „X\$*”
 - General wild-card query: for „X*Y” we look up „Y\$X*”
- **Q:** How would you handle the query „X*Y*Z” with the permuterm index?
 - **A:** Here we have no option but to fire two lookups into the index
 1. Retrieve the postings for „X*Z” (by looking up „Z\$X”)
 2. Retrieve the posting list for the query „*Y*” (by looking for „\$Y”)
 3. Intersect the two retrieved lists of terms

Character indexes

22

- Next idea:
 - How about we index all **character n-grams** (sequences of n characters) instead of whole terms?
 - We surround all terms with term-boundary symbols (\$) and create lists of all sequences of **n** consecutive character within terms
 - Example: „**Frodo and Sam fought the orcs**” (stopwords removed; lemmatized)
 - Terms: \$frodo\$, \$sam\$, \$fight\$, \$orc\$
 - Char. 3-grams: \$fr, fro, rod, odo, do\$; \$sa, sam, am\$; \$fi, fig, igh, ght, ht\$; \$or, orc, rc\$
- We need to keep the **second inverted index**
 - For each character n-gram maintain the list of vocabulary terms that contain it
 - E.g., „\$fr” -> [„freak”, „freedom”, ..., „frodo”, „frozen”]
„sam” -> [„asamoah”, „balsam”, „disambiguate”, ..., „sam”, „subsample”]

Wild-card queries and character indexes

23

- Query for character n-grams and merge results (AND operator)
- Example: query „mon*” and 3-gram character indexing
 - Query is transformed into: „\$m” AND „mo” AND „on”
 - **Q:** What might be the issue with this transformation?
 - **A:** Conjunction of character 2-grams might yield false positives
 - For example: moon, motivation, moderation, etc.
 - Compare this issue with the false positives of biword index from Lecture #2
 - Retrieved terms **must be post-filtered against the query** to eliminate false positives
 - Term contains „mon”?
- Resulting terms are then looked up in the term-document inverted index

Character indexes

24

- Comparison with permuterm index
 - **Advantage:** space efficient (less space needed than for permuterm index)
 - **Shortcoming:** slower than using permuterm index
 - A Boolean query (and term-level merges) needs to be performed for every query term
- Wild-card queries in general
 - Often not supported by Web search engines (not at the character level anyways)
 - Found in some desktop or library search systems
 - Wild-cards are **conceptually troubling** as well
 - User must know what they don't know (i.e., where to put the asterisk)
 - Used a lot in domain-specialized early Boolean retrieval systems – in part as a **term normalization** technique (before stemming or lemmatization were widespread)
 - If we have several options in mind, we can just run several concrete queries

Outline

25

- Recap of Lecture #2
- Data structures for inverted index
- Wild-card queries
- Spelling correction

Spelling correction

26

- Primary use-cases for spelling correction
 1. Correcting documents during indexing
 2. Correcting user queries on-the-fly

- Two flavors of spelling correction
 1. Isolated words
 - Check each word on its own for errors in spelling
 - Will **not** catch typos that result in another valid word
 - E.g., „from” → „form”
 2. Context-sensitive spelling correction
 - Correctness evaluated by looking at surrounding words as well
 - E.g., „Frodo went form Gondor to Mordor”

Document correction

27

- Correction should occur prior to indexing
 - Aiming to have only valid terms in the vocabulary
 - Smaller vocabulary, i.e., the term dictionary contains fewer entries
- We do **not** change the original documents
 - Just perform correction when normalizing terms before indexing
- Common types of errors for certain types of documents
 1. OCR-ed documents – „rn” vs. „m”, „O” vs. „D”
 2. Digitally-born documents often have QWERTY keyboard typos – errors from close keys – „O” vs. „I”, „A” vs. „S”, etc.

Query correction

28

- Primary focus is on correcting errors from queries
 - **Q:** Failing to fix errors in queries has more serious consequences than omitting to fix errors in documents. Why?
- With respect to user interface, we have two options
 1. Silently retrieving documents according to the corrected query
 2. Return several suggested „corrected” query alternatives to the user
 - „Did you mean?” option

Isolated word correction

29

- **The idea:** using **reference lexicon** of correct spellings (i.e., lexicon of valid terms)
- Two approaches for obtaining a reference lexicon
 1. Existing lexicons like
 - Standard wide-coverage lexicon of a language (e.g., Webster's English dictionary)
 - Domain-specific lexicons (e.g., lexicon of legal terms)
 2. Lexicon built from large corpora
 - E.g., all the words on the web or in Wikipedia
 - **Q:** Do we want to keep absolutely all terms from large corpora?

Isolated word correction

30

- Given a reference lexicon and the query term (a character sequence from the query), we do the following:
 1. Check if the query term Q is in the reference lexicon
 2. If the term Q is not in the reference lexicon, find the entry Q' from the lexicon that is „closest” to the query term Q

- How do we define „closest”?
 - We need some similarity/distance measure
 - We will examine several options
 1. Edit distance (also known as Levenshtein distance)
 2. Weighted edit distance
 3. Character n-gram overlap

Spelling correction – edit distance

31

- **Edit distance** between two strings S and S' is the **minimal number of operations** required to transform one string into the other
 - What are the „operations“?
- We typically consider operations at the character level
 - Character **insertion** („frod“ \rightarrow „frodo“)
 - Character **deletion** („frpodo“ \rightarrow „frodo“)
 - Character **replacement** („frido“ \rightarrow „frodo“)
 - Less often: **transposition of adjacent characters** („fordo“ \rightarrow „frodo“)
 - Transposition equals „deletion“ + „insertion“?
 - **Q:** Why introducing it as a separate operation?
- **Levenshtein distance**: counts insertions, deletions and replacements
- **Damerau-Levenshtein distance**: additionally counts transpositions as a single operation
- Algorithm based on **dynamic programming**

Dynamic programming

32

- For detailed explanation of dynamic programming see [Cormen, Leiserson, Rivest, and Stein. „Introduction to Algorithms“](#)
- **Optimal substructure**: the optimal solution of the problem contains within itself the subsolutions, i.e., the optimal solutions to subproblems
- **Overlapping subsolutions**: we can recycle subsolutions – i.e., avoiding repeating the computation for the same subproblems over and over again
- **Q:** What would be a „subproblem“ for the edit distance computation?
 - **A:** the edit distance between two prefixes of input strings
- **Q:** Do we have many subproblem repetition for edit distance?
 - **A:** most distances between same pair of prefixes are needed 3 times (as a subproblem of computing distance for insertion, deletion, and substitution)

Levenshtein distance

33

- Let a and b be two strings between which we measure edit distance (with $|a|$ and $|b|$ being their respective lengths):
- Mathematically, the Levenshtein distance $lev_{a,b}(|a|, |b|)$ is computed as follows:

$$lev_{a,b}(i, j) = \begin{cases} \max(i, j) & \text{if } \min(i, j) = 0, \\ \min \begin{cases} lev_{a,b}(i-1, j) + 1 \\ lev_{a,b}(i, j-1) + 1 \\ lev_{a,b}(i-1, j-1) + 1_{(a_i \neq b_j)} \end{cases} & \text{otherwise.} \end{cases}$$

- Where $1_{(a_i \neq b_j)}$ is the indicator function equal to 0 if $a_i = b_j$ and 1 otherwise
- Once we compute $lev_{a,b}(i, j)$ for some pair (i, j) we store it in memory so we don't compute it again when needed in another recursive thread
- Directly implementing this formula requires **recursion**

Example – Levenshtein recursively

34

- For the example, we will follow only one thread of recursion (first subproblem)
- „sany” vs. „sam”
 - $\min(\text{lev}(\text{„san”}, \text{„sam”}) + 1, \text{lev}(\text{„sany”}, \text{„sa”}) + 1, \text{lev}(\text{„san”}, \text{„sa”}) + 1)$
- „san” vs. „sam”
 - $\min(\text{lev}(\text{„sa”}, \text{„sam”}) + 1, \text{lev}(\text{„san”}, \text{„sa”}) + 1, \text{lev}(\text{„sa”}, \text{„sa”}) + 1)$
- „sa” vs. „sam”
 - $\min(\text{lev}(\text{„s”}, \text{„sam”}) + 1, \text{lev}(\text{„sa”}, \text{„sa”}) + 1, \text{lev}(\text{„s”}, \text{„sa”}) + 1)$
- „s” vs. „sam”
 - $\min(\text{lev}(\text{„”}, \text{„sam”}) + 1, \text{lev}(\text{„s”}, \text{„sa”}) + 1, \text{lev}(\text{„”}, \text{„sa”}) + 1)$
- „” vs. „sam”
 - return 3

Levenshtein distance – non-recursive version

35

- We can **avoid the recursion** if we start from the recursive algorithm's end condition – return **$\max(i, j)$** if **$\min(i, j) = 0$**
- Then compute the edit distances of larger prefixes from smaller prefixes

LEVENSHTEINDISTANCE(s_1, s_2)

```
1  for  $i \leftarrow 0$  to  $|s_1|$ 
2  do  $m[i, 0] = i$ 
3  for  $j \leftarrow 0$  to  $|s_2|$ 
4  do  $m[0, j] = j$ 
5  for  $i \leftarrow 1$  to  $|s_1|$ 
6  do for  $j \leftarrow 1$  to  $|s_2|$ 
7      do if  $s_1[i] = s_2[j]$ 
8          then  $m[i, j] = \min\{m[i-1, j]+1, m[i, j-1]+1, m[i-1, j-1]\}$ 
9          else  $m[i, j] = \min\{m[i-1, j]+1, m[i, j-1]+1, m[i-1, j-1]+1\}$ 
10 return  $m[|s_1|, |s_2|]$ 
```

Example – Levenshtein non-recursively

36

	–	s	a	m
–	0	1	2	3
s	1	0	1	2
a	2	1	0	1
n	3	2	1	1
y	4	3	2	2

Damerau-Levenshtein distance

37

- Standard edit distance counts **transposition of adjacent characters** as two edits
 - E.g., „frodo” vs. „fordo”
 - two character replacements: „r” -> „o” in position 2 and „o” -> „r” in position 3
- However, transposing adjacent characters is usually a **single typing error**
 - **Damerau-Levenshtein distance** introduces **transposition** as the fourth atomic distance operation
 - **Q:** How would you integrate transposition as a single distance operation into the edit distance algorithm?
 - **A:** $d(i,j)$ additionally needs to consider $d(i-2, j-2) + 1(a_{i-1} = b_j \ \& \ a_i = b_{j-1})$ when looking the edit distances of prefixes

Weighted edit distance

38

- Sometimes we want to assign smaller distance to common errors
 - The weight of an operation (deletion, insertion, replacement, transposition) depends on the character(s) involved
- Motivation: better capture common OCR or typing errors
 - E.g., On a QWERTY keyboard, letter „m” is much more likely to be mis-typed as „n” than as „q”
 - Thus, the replacement operation „m” -> „n” should be assigned smaller edit distance than „m” -> „q”
- Additional input required
 - Data structure (e.g., weight matrix) containing operation weights for (combinations of) characters
- **Q:** How to integrate weighting into the edit distance algorithm based on dynamic programming?

Using edit distances

39

- Given a (misspelled) query we need to find the closest dictionary term
- **Q:** How do we know (or assume) that the query is misspelled in the first place?
 - **A:** We don't find the query term in the vocabulary dictionary
 - With this strategy, we cannot capture typos like „from” -> „form”
- Finding closest dictionary term
 - Compute edit distance between the query term and each of the dictionary terms?
 - **Too slow** (the dictionaries are usually rather large)
 - We need to somehow **pre-filter** the „more promising” dictionary entries

N-gram index for spelling correction

40

- **Idea:** use the **character n-gram index** to pre-filter dictionary candidates
 1. Enumerate all character n-grams in the query string
 - E.g., 3-grams in „frodso” -> „fro”, „rod”, „ods”, „dso”
 2. Retrieve all vocabulary terms containing any of the obtained character n-grams
 - Using the **inverted index of character n-grams**
 3. Treshold the obtained list of candidates on the number or percentage of matching character n-grams
 4. Compute the edit distances between the query term and the remaining dictionary candidates
 5. Select the candidate with the smallest edit distance as the correction

Character n-gram overlap

41

- Can be used as
 - A measure for pre-filtering candidates in order to reduce the number of edit distance computation
 - As a self-standing distance measure, alternative to Levenshtein distance
- Example
 - Suppose the query is „fpodo bigginss” and the text is „frodo baggins” and we are computing the overlap in character 3-grams
 - {„fpo”, „pod”, „odo”, „big”, „igg”, „ggi”, „ins”, „nss”} vs. {„fro”, „rod”, „odo”, „bag”, „agg”, „ggi”, „ins”}
 - We have 3 matching 3-grams: „odo”, „ggi”, and „ins”
 - That’s 3 out of 8 for the query and 3 out of 7 for the text
- **Q:** What should we take as measure of proximity/distance?
 - Is raw count of matching n-grams good choice?

Character n-gram overlap

42

- Raw count of matching character n-grams **is not a good choice**
 - Does **not account** for the **length of terms** in comparison
 - Two distinct but long terms may have a large raw count of matching n-grams
 - E.g., „collation” and „collaboration” have 5 matching 3-grams
 - We need to normalize the score with the length of terms

- **Jaccard coefficient** – a commonly used measure of set overlap

$$|X \cap Y| / |X \cup Y|$$

- Simple alternative: averaged length-normalized overlap

$$0.5 \cdot (|X \cap Y| / |X| + |X \cap Y| / |Y|)$$

Context-sensitive spelling correction

43

- Example:
 - Suppose the text is „Frodo fled from Mordor back to Gondor”
 - Suppose the query is „fled form Gondor”
- To identify the misspelling „form” -> „from” we need to take into account the context, i.e., surrounding words
- Context-sensitive error correction steps
 1. For each term of the query, retrieve dictionary terms that are sufficiently close
 - „fled” -> {„fled”, „flew”, „flea”}; „form” -> {„form”, „from”}; „gondor” -> {„gondor”}
 2. Combine all possibilities (i.e., all combinations of candidates for each term)
 - „fled form gondor”, „fled from gondor”, „flew form gondor”, „flew from gondor”, „flea form gondor”, „flea from gondor”,
 3. Rank the possibilities according to some criteria

Context-sensitive spelling correction

44

- **Hit-based** spelling correction
 - Rank the candidate combinations according to the number of hits (no. documents that contain those combinations)
 - Return the candidate with the largest number of hits
- **Log-based** spelling correction
 - Rank the candidates according to the number of appearances in the query logs (i.e., the number of times the same query was posed before)
 - Useful only if you have a lot of users who fire a lot of queries
- **Probabilistic** spelling correction (e.g., based on **language modeling**)
 - Ranking according to probabilities of term sequences
 - E.g., $P(\text{„fled form gondor”}) = P(\text{„fled”}) * P(\text{„form”} \mid \text{„fled”}) * P(\text{„gondor”} \mid \text{„form”})$

Now you...

45

- Know what data structures you can use for implementing inverted index
- Understand the pros and cons of hashtables and trees
- Know how to handle wildcard queries
- Are familiar with methods for handling spelling errors and typos in IR